

How Much Diversity Does a Red-Teamer Need?

Minimum training diversity for transferable multi-step indirect prompt injection

Yang Gao • Fortinet • 2026-09 • Project overview

In July 2026 OpenAI showed that an RL-trained red-teamer generalizes to held-out environments and defenders — using compute on the scale of a GPT-5 post-training run. Other work at small scale reports the opposite: that RL attackers memorize their training targets. Both can be true. Nobody has measured the curve in between. **This project measures it**, on an attack surface where the payload itself is a learned decision: how many innocuous fragments to split an injection into, and which points along the agent's trajectory to place them at.

0. The question



Fig 0 — Two endpoints are published. The curve between them is not.

A big lab does not need this answer — it can simply scale until the problem goes away. **A small, honest measurement of where the knee is** costs orders of magnitude less than the run that made it necessary, and is useful to everyone who cannot afford a GPT-5-scale one.

1. The problem




Fig 1 — Indirect prompt injection: the harm is an action, not a sentence.

The harm is not that a model said something offensive. The harm is that an agent with real tool access **took an action** — it sent the email, transferred the money, pushed the commit.

2. What exists today

Work	What it does	What it does not decide
GPT-Red 2607.26115	RL self-play attacker vs. a population of defenders. "The largest published safety RL run to date" — compute comparable to post-training GPT-5. Generalizes to held-out environments and defenders.	No decomposition, no placement choice. Its attacker iteratively rewrites one payload against a stateful defender tool. Not released — no weights, code or environments.
AgentDojo templates	Fixed hand-written strings	Everything — they are constants
StepJack 2608.06477	Splits the goal into k fragments	k is a CLI flag; placement follows page order; pipeline is explicitly open-loop
AutoInject 2602.05746	Black-box RL learns a suffix	Single injection point — learns what to write, not where
PISmith 2603.13026	GRPO attacker vs. defenses	One complete payload; position fixed by a concatenation operator
PIMiner 2608.05108	Agentic strategy library	No decomposition at all; injection site given by the benchmark

The pattern is consistent, and it survives GPT-Red. **The field has automated what to write — at every scale from a single GPU to a GPT-5-sized RL run. Nobody has automated how many pieces, and where they go.**

That gap is not cosmetic. StepJack's own numbers show why: multi-step decomposition helped on only **3 of 6** target agents, lifting average attack success from **31.3% to 36.9%**. That variance is the signature of a one-size-fits-all decomposition meeting targets that need different ones.

3. What this work adds




Fig 2 — Single payload vs. learned decomposition and placement.

Decision	Plain meaning
Adaptive depth	How many pieces to cut the goal into — chosen per victim, not set by a flag
Learned placement	Which of the slots the agent will actually read should carry a fragment
Intra-trajectory adaptation	Mid-attack, look at which tools the victim already called, then adjust the unused slots

A property worth stating: the reward is verifiable

AgentDojo returns a hard 0/1 : did the agent execute the attacker's intended tool call. **No judge model, no reward model, no room for reward hacking through a learned scorer.** That is rare in security ML, where most signals are one model's opinion of another's output.

4. System




Fig 3 — Role reversal: the attacker is the policy; the victim agent is environment dynamics.

5. Evidence so far




Fig 4 — Qwen3-4B + LoRA, GRPO, single injection point. Two independent runs on different machines and torch versions: template attack 57.8% / 62.5%, trained policy 100% / 100% (n=64 each); held-out tasks 2 and 3: 31.3% / 37.5% (n=16) and 31.3% / 40.6% (n=32). Held-out task 1 is excluded from these numbers until it is re-tested: the trained policy scored 0/48 on it across both runs, every attempt reaching the target step and failing there, and the template attack scored 1/16 — a pattern we have not yet explained.

The 100% establishes that the signal is learnable. **The 31% and 41% on held-out tasks 2 and 3 are not a failure — they are the leftmost point of Fig 0**, measured at the minimum possible diversity: one task, one victim.

A second, smaller result from the same run is worth stating because it is about the same variable. The attacker prompt was originally written in Chinese while the AgentDojo environment, victim and injection goal are all English. Holding model, temperature, sample count and task pair fixed and changing only the prompt language: ASR 0.00 → 6.25% and within-group reward variance 0.000 → 0.059. Below that variance GRPO has no gradient at all. **A mismatch between the attacker's output distribution and the environment did not degrade learning — it zeroed it.**

The diagnosis is independently corroborated. GPT-Red reports the same failure at their scale (§ 5.1): "we ran self-play with only a **single defender** ... [this] often resulted in attackers [that] would **mode collapse**." **The cause is training-distribution width, not the RL method.**

6. Where the project stands




Fig 5 — Four rungs. Each is a complete, checkpointable GRPO run; each adds exactly one unseen thing.

L2 is now designed as the sweep itself: train at increasing diversity — tasks, then suites, then victim models — and plot held-out retention at each setting. That produces Fig 0's missing curve as the paper's primary result, with learned decomposition and placement as the mechanism being swept.

7. Is there a future in this?

Arguments for

- **The framing survives the biggest result in the field.** GPT-Red automates what to write at enormous scale and still does not decide decomposition or placement.
- **The threat is growing.** Agent deployments with tool access keep expanding; recent benchmarks report direct injection succeeding against GPT-5- and Gemini-class agents over **79%** of the time.
- **The gap is unclaimed.** Verified by reading the three most recent competing systems' papers and repositories — not their abstracts.
- **Clean RL problem.** Verifiable binary reward, a real environment, a well-defined action space.
- **The defense side needs it.** A benchmark where every defense scores near-perfect is not measuring deployability.

Arguments for caution

- **Being read as "GPT-Red, but smaller."** This is now the single largest risk. The defence is to lead with the measurement question, not with the attacker.
- **A well-resourced group is in the same territory.** One lab shipped the benchmark, the RL attacker and the agentic attacker inside six months. Speed to arXiv matters.
- **Generalization was the exposed flank — less so now.** A competing paper's headline claim is that RL attackers transfer poorly. GPT-Red refutes that as a general statement, which converts the open question from "can it?" to "how much is needed?" — the question this work asks.
- **The benchmark is aging.** No package release in nearly a year; its defense side is close to saturated.
- **This is a 7-8 week window.** The area moves in weeks, not months.

Assessment. An academic-scale project cannot out-scale a frontier lab, and should not try. What it can do is answer a question that scaling makes uninteresting to ask: **where is the knee in the diversity curve?** The two endpoints are published and the middle is empty. The contribution is a measurement plus an open, reproducible artifact — precisely the two things the largest result in the field does not provide.

8. Budget

The curve in Fig 0 is the deliverable, and its price is part of the argument. GPT-Red answered the question at a scale only a frontier lab can reach. **This asks for orders of magnitude less to measure the part that scale skipped** — where the curve turns.

Item	Amount	Basis
RunPod H200-141GB 120 GPU-h @ \$3.59/hr	\$431	Built bottom-up from measured numbers. Floor, measurements only: 24 runs (8 diversity points $\times$ 3 seeds) $\times$ 300 steps $\times$ 27.1 s/step (measured over 100 steps) = 54 GPU-h. $\times$ 1.2 for larger victims: only the rollout share of a step ($\sim$ 10 of 27 s) grows with victim size, and only at the points that use one. $\times$ 1.5 for re-runs: L1 needed three attempts to produce a clean run. Plus $\sim$ 8 GPU-h to reproduce baselines in-environment. $54 \times 1.2 \times 1.5 + 8 \approx 105$ GPU-h, requested as 120 . Priced at H200 because one card must hold the 4B attacker in training alongside a 31B victim; step times were measured on a slower A100, so this is conservative.
Frontier-model API credits	\$2,000	Held-out transfer evaluation against five frontier targets at $n = 200$ per condition, plus faithful reproduction of the PIMiner baseline, which requires a frontier model as its attacker backbone. No API calls occur inside the training loop.
Contingency 20% of API	\$400	Compute re-runs are already inside the $\times$ 1.5 above, so contingency is not applied to compute a second time.
RunPod network volume 1 TB $\times$ 5 months @ \$0.07/GB/mo	\$350	Persists weights, checkpoints and results between pods. Billed continuously, including while no pod is running.
Total	\$3,181	

Cost discipline is already in place, and it was bought the hard way. An earlier pod idled 62 hours at 0% GPU utilization after a run early-stopped and nothing was watching — about \$98 wasted. Shutdown is now a mechanism rather than a line in a runbook: a watchdog on the pod stops it after 20 minutes with no training process and GPU below 5%, with a hard lifetime cap behind that. **Every number in this document comes from an archived JSON under docs/research/ ; figures observed but not archived are not quoted.**

References: [GPT-Red arXiv:2607.26115](#) • [StepJack arXiv:2608.06477](#) • [AutoInject arXiv:2602.05746](#) • [PISmith arXiv:2603.13026](#) • [PIMiner arXiv:2608.05108](#) • [AgentDojo arXiv:2406.13352](#)